

Table 1, every version of it

A design dossier for the breakdown-radius paper

José Lucas de Melo Costa · 16 September 2026

Internal document. Author lock 18 September, paper 25 September.

1 What this is, and what it was built from

You asked for every Table 1 the results can support, with the claims behind them, their significance and their originality, and with the comparison to the literature made visible. This document holds thirteen candidates, all of them typeset with the numbers that exist today, plus the claim set they are supposed to verify and the scorecard that separates them.

Almost nothing here was invented this morning. Between 7 and 9 September five workflow runs designed, attacked and partly measured this table and the quantity it reports: 210 agents and 26.2 M context tokens over about nine and a half hours. Between 11 and 13 September the measurement runs filled the cells those designs had left empty. This dossier reads the reports and the result files, keeps what survived its own adversaries, and adds the candidates that only became possible once the measurements existed. Where a number appears in a table below, the file it was read from is printed under the float; where a number was computed for this document, the query that computed it is named.

date	design or measurement	agents	context	what it left for Table 1
07/09	table ideation	60	6.0 M	28 candidates, 21 kept; no discrimination table survives
07/09	table selection	15	2.2 M	the four findings F1–F4 that every later design answers
07/09	main-table design	43	5.4 M	procedures as rows, one estimator; the screen identity
07–08/09	radius redefinition	67	9.1 M	retraction is reversal; the null radius; the far control
09/09	archetype ranking	25	3.4 M	ten archetypes attacked; the repairs W1–W7
09/09	apostila Part IV	–	–	one calibrated metric; columns by number of wrong claims
11–13/09	the protocol run	8 arms × 528 rows		the certificate ledger, 27 networks
13/09	M1, the knowledge interval	2,749 rows		the price table, calibrated, with bands
13/09	M1b, the policy rows	2,212 rows		the refusal rate and the tie against the free screen

1.1 The situation this table has to survive

Chris’s draft arrived at 01:59 on 16 September: a complete paper of nineteen pages, eleven of them main text, titled *A Diagnostic Radius for Knowledge-Induced Causal Effect Estimation*. His Table 1 is a rank-correlation table, τ_b between each instrument and a survival area, by stratum, over three million corrupted knowledge states on 1,978 synthetic instances. The two experiments his draft names as still owed are the two we ran on 11–13 September: the ranking evaluation on the real instances, and the elicitation arm with a frozen instance frame, a matched control and an audit opened after the rows were frozen.

The merge recommendation on record is his skeleton with our results section. That decision changes what Table 1 can be, in one specific way: the table has to serve a paper whose theory section is his geometry and whose evaluation is our measurement, and it has to hold the baseline his table does not have.

A full review of that draft was written earlier today and this dossier is its companion rather than a second opinion. Two of its findings are carried here. The first is about the theory and it is in C2: the order on this space is graded because the skeleton is frozen, which is consistent with the result in Taeb et al. that the general poset is not. The second is about the free baseline and it is in C5 and T8: the defensible claim is that the radius says which depth was needed, and the coverage is won there.

1.2 The three requirements, and the constraints already paid for

Your brief sets three requirements. A reviewer who has read only the abstract and the introduction must be able to read the table (R1); the table must verify the claims we actually make (R2); and it must show that we compared ourselves to the literature (R3).

Six further constraints are not preferences. Each replaces something that was measured and found to fail, and every candidate below is scored against them.

	the constraint	what it replaces, and where that was measured
R4	One metric, not two free columns	A coverage column and a width column let any procedure shelter in one of them. Discarding the knowledge covers 1.000 in every regime and passes any coverage gate, so the argument has to happen in the width column, which was empty. The repair is to calibrate every procedure to 95 % coverage and then measure the width it needed, relative to the width of discarding the knowledge.
R5	One population per panel, and every stratum printed with its n	Three columns drawn from three non-nested populations, the defect the ideation audit found in the coverage family. A thin cell is printed with its count rather than suppressed.
R6	A row that loses, in the column where it must lose	Insurance costs something when nothing is wrong. A table whose favoured row wins every column is a design error and reads as one.
R7	No discrimination column	The expected rank of a reversed claim is $1 + (K - 1)(1 - \text{AUC})$, so a questions-to-repair column restates an area under a curve that already ties. No table in which the radius competes on discrimination survived the ideation run.
R8	No sentence whose subject is a language model	Elicited knowledge is one error process in the methodology, never the headline, the title or Table 1.
R9	Nothing internal in the float	No script names, run identifiers or hashes in a caption. This document breaks that rule on purpose, because it is not the paper; the lines marked as draft captions are written as they would be printed.

2 The claims, and what each is worth

Eleven claims are live. Three of them are the paper: the object (C1), the unit (C3) and the price (C4). Two are honest negatives that a reviewer will find if we do not print them (C5, C6), and one is a scope limit that reads badly and has to be printed anyway (C8). The rest support or bound them. Each is written below in the form it is licensed in, which is narrower than the form it is tempting to write.

Every claim carries a line naming what would kill it. Where the register records that a novelty search was never run, the claim is marked as blocked rather than dressed as verified.

C1 The object: nobody prices a structural assumption

Sensitivity analysis in causal inference is built for unmeasured confounding and for continuous relaxations of identifying assumptions. The assumption that background knowledge carries is discrete and structural, and it is not priced the same way. An analyst who asserts orientations, propagates them and reads an adjustment set off the result has no number telling them how much of the reported effect rests on those assertions. The assertions are checked for consistency with the data and never for truth, so a claim that is consistent and false passes every check that current practice has.

rests on	the framing, not a measurement; the measurements below are what makes it a paper
significance	this is the one place where unverifiable human input enters an otherwise data-driven pipeline, and the estimate at the end of the pipeline inherits it silently
nearest prior art	Taeb, Guo and Henckel 2025 (arXiv:2511.10625), model-oriented graph distances over posets with an adjustment-set robustness application. This is the foundation, not the competitor, and the paper says so. Guo and Perković (arXiv:2010.08611) enumerate adjustment sets in an MPDAG and never perturb the knowledge. Masten and Poirier, <i>Quantitative Economics</i> 11(1):41–111, 2020, define breakdown frontiers with inference, in econometrics and not on graphs: prior art for reporting a breakdown value with a band, outside machine-learning venues. Rosenbaum, Cinelli and Hazlett, Ding and VanderWeele all price a continuous confounding parameter
falsifier	a tolerance radius centred on the analyst’s own graph inside Taeb et al. §6, which nobody on this side has read in full
carried by	T3, the report on named analyses, written in the Rosenbaum form

Blocked, and it blocks the abstract. The register entry for this claim carries `verified_NOT_RUN`: no novelty search has ever been run in this wording. The three sweeps on record were pointed at latent-causal representation learning and at Taeb specifically, not at the sensitivity-analysis lineage the claim excludes, and the framing rests on the statement that Guo and Perković never perturb the knowledge, from a paper nobody has read in full. Under the submission rule a novelty claim verified more than thirty days ago cannot be reused, and this one was never verified at all. The search has to run before the abstract goes in on 18 September, and C1 has to be narrowed to whatever it returns.

C2 The quantity is computable from the analyst’s own inputs

The breakdown radius is the distance from the analyst’s graph to the nearest knowledge state at which the committed adjustment set stops being valid. It is computed from the estimated CPDAG, the claims, the treatment, the outcome and the committed set. The data-generating graph is on no path of the computation; it is used to draw samples and, afterwards, to check whether the certificate held.

rests on	failure is upward-closed, the space is a join-semilattice with the analyst’s graph as its maximum, the nearest failure is reached by retracting only, and cover generation is set arithmetic. Our own contribution to this is the audit: 5,000 sampled pairs over 469 rows across the elicited, truthful and oracle conditions, zero violations, one-sided 95 % upper bound on the violation rate of 0.06 %
significance	the reachability result rests on an anti-exchange property that is verified and not proved, with a one-sided error that overstates robustness. The audit is the direct evidence for it, and it is a number where the draft currently has a caveat
the positioning	this morning’s review of the draft found the sharpest version of the point. Taeb et al. prove in their Proposition 9 that the poset of MPDAGs is not graded for four or more vertices and not semimodular, with no characterisation of the covering relation; the draft’s appendix defines a rank function and concludes that its order is graded. Both are correct, because their counterexample varies the skeleton and this space holds it fixed. Saying so converts the largest liability in the theory section into the clearest statement of what fixing the skeleton buys
nearest prior art	the order and the distance are Taeb et al.; the specialisation to knowledge states and the validity predicate are ours and Chris’s
falsifier	one counterexample to upward closure halts the argument. None found in 5,000 draws
carried by	T5 makes it legible, T2 makes it safe

C3 The unit: a certificate in claims is safe, a certificate in graph revisions is not

This is the strongest thing the measurement produced and it is the one place where our results change what the draft reports. A certificate denominated in the analyst’s own claims never certified a set that was invalid at the truth, on any row of any condition. The same certificate denominated in elementary graph revisions did.

rests on	3,014 certifiable rows under eleven knowledge conditions over 27 networks: the claim unit over-certifies on 0, the revision unit on 39, a free rule read off the graph on 15, and the two naive counts on 729 and 791
significance	it decides the unit the paper reports in. The draft reports the radius in elementary revisions and treats the claim count as a practitioner-facing refinement; the ledger says the refinement is the safe one and the revision unit is not. It is also the reason the certificate has to be redenominated rather than rescaled
the argument	structural, and the paper must not present it as statistical. If the committed set is invalid at the truth and the condition asserted w false claims, retracting those w claims leaves a graph the truth extends, on which the criterion fails, so the claim radius is at most w whenever the search reaches depth w . The zero is what a sound criterion and an audit asking the same question must return, and it was returned once the audit did
nearest prior art	none. No accepted paper in the record we searched computes a breakdown radius or a null radius, so there is no precedent for the comparison and no borrowed template; the comparison set is the instruments an analyst already has
falsifier	one row on which the claim certificate certifies a set invalid at the truth, under a correct generalised adjustment criterion. Two such rows existed before the audit oracle was fixed, and both were the oracle
carried by	T2, and it is the reason T2 is the recommendation

C4 The price: what the interval costs, and what it buys back

When the claims are right, the interval that tolerates one revision costs about twice the width of the report as committed. When one claim is wrong, the report as committed covers the true effect on six problems in ten, and the interval covers on ninety-eight of a hundred at 85 % of the width of discarding the knowledge altogether. Two wrong claims need two retractions: the interval at one revision fails and the interval at two pays.

rests on	the calibrated width table, 460 informative rows over 23 networks, with a network cluster bootstrap; predictions P1, P2, P4 and P8 supported in four cells of four, P3 partial, P6 falsified, P7 supported
significance	this is the sentence a practitioner acts on, and the only one in the paper that gives a price in the currency they report in
nearest prior art	the class-wide interval of Guo and Perković is the procedure that discards the knowledge, and it is the denominator of every cell. The omitted-variable bound of Cinelli and Hazlett is the sensitivity row, and it is an oracle here
falsifier	the interval at one revision failing to cover at one wrong claim, or costing more than discarding the knowledge. Neither happens in any cell
what it must not become	a continuity bound. Damage beyond the radius is heavy-tailed, so “small knowledge error implies small bias” is false and our own data falsifies it
carried by	T1

C5 The tie: the validity half computes what a free screen computes

A reviewer who knows the area will ask why an analyst needs a radius when they could leave one claim out, close the graph again and look at whether the set changed. The answer is that at one revision they could not tell the difference, because the two are the same policy, and at two revisions they tie exactly when the screen is also run at two.

rests on	the stop rule and the leave-one-out screen decide differently on 2 of 162 rows at one wrong claim, both of the class the argument predicts. At two wrong claims the stop rule beats the leave-one-out screen by 0.305 [0.182, 0.410] in coverage at the population and by 0.256 [0.149, 0.346] at a sample of one thousand, and ties the leave-two-out screen at 0.000 [0.000, 0.000]. The separation the design asked for exists only because that baseline stops at depth one
the sentence this licenses	a screen answers yes or no at whatever depth the analyst happened to choose. The radius says which depth was needed, and the coverage is won there. Quote one of the two numbers above with its sample size attached; they are the same contrast measured twice
significance	this was pre-declared on 7 September, with the sentence for the null written before the measurement, which is the reason it can be printed as a property rather than conceded as a defeat. It is also the baseline the draft’s ranking table does not have, and the single most valuable thing our side can hand that table
nearest prior art	the screen is not in the literature; it is the thing a competent reader invents in ten seconds, which is worse
the licensed sentence	the validity half of the report computes what a leave- k -out re-closure screen computes, and the paper states this as a property of the measure. What the radius adds is the depth, the named witness and the certificate that no failure lies nearer
carried by	T8, and one row of T1

C6 The guarantee is expensive exactly where the knowledge matters

On the problems where the knowledge carries identification, retracting a single true claim breaks the committed set on nine problems in ten. The refusal rate is 0.910 on the informative rows with a committed set, zero on the committed rows the class already identifies, and 0.639 over all committed rows. The denominator decides the number, so the paper states the denominator.

rests on	134 informative committed rows over 20 networks in the semi-synthetic panel, 7 rows over 3 networks in the fitted panel; the prediction that the rate would stay below one half was falsified
significance	this is the finding that stops the paper being a method paper. Knowledge buys identifiability rather than a better set, and the certificate mostly says the conclusion is not robust. That is a real diagnostic result and it reads as one only if we print it ourselves
falsifier	a population on which the refusal rate is low and the certificate still detects; the informative stratum is not it
carried by	T8

C7 Locality holds in the unit the analyst asserts, and not in hops

Hop distance predicts the depth when the binding claim is incident to the treatment, and overstates it otherwise. On the primary condition’s 157 exact rows the hop reading equals the claim depth on 135, exceeds it by one on 19 and by two on 2, and falls below on 1. Reading above the claim depth is what over-certification is.

rests on	the agreement counts above, computed for this document; and the matched-budget negative control, where retracting a claim far from the query is not inert on the semi-synthetic panel (coverage +0.044 at one wrong claim) and is inert on the fitted panel
significance	locality pays in cost and in what is worth eliciting, and does not pay in width
what it must not become	“knowledge at distance at least one cannot harm”. The 0/791 figure was struck in August for having no provenance and is never repeated
carried by	T13, or the compact locality section

C8 Applicability, stated as a limit

Of the nine author-declared queries in the corpus that can be certified, eight are unbroken by any retraction down to depth three. Of the sampled population, the committed set breaks at the first retraction on almost every row. The mechanism is the sampling rule: a pair is admitted only when some retraction changes validity, so a robust query is rejected before it is measured.

rests on	the observables-only frame, 80,132 candidate pairs over 33 networks, against 543 under a gate that reads the true graph. Eight networks yield a frame where the old gate yielded no instance at all
significance	the measured population is close to the complement of the queries the source papers exist to answer. A reviewer who notices this before we state it has found our worst objection; a reviewer who reads it in our own funnel has read a scope statement
carried by	T6, and it belongs in the funnel rather than in Table 1

C9 Scope: the graph is the oracle’s

The licensed phrasing is that the radius is computable from the estimated CPDAG, the claims, the covariance and the sample size. It is never that the conclusion is robust under an estimated CPDAG. On the learned-graph panel the true graph extends the analyst’s graph on none of the rows, and both instruments over-certify on 14 of 38.

C10 The supplier is an error process, and never the headline

Measured commission and validity by condition, with a matched chance control: the model beats chance in commission by 0.143 [0.006, 0.287] and is not demonstrated to beat it in validity, 0.082 [−0.020, 0.176]. Scale buys claims and accuracy and not validity: the share of committed sets invalid at the truth is 0.298 at the largest supplier against 0.288 at the smallest. This is methodology. It is ruled out of the Table 1 slot.

C11 Ranking: the radius orders instances by survival better than the standard distances

The radius ranks instances by whether the committed set survives corrupted knowledge better than structural Hamming distance and than the number of claims, in most strata. Asserting more knowledge is associated with less robustness. Structural Hamming distance is a poor predictor: knowledge can be entirely correct and the set still break at the first subsequent error. One cell is a confirmed null and two are unstable at 200 draws.

Not claimed, in any draft. No continuity bound. No probability reading of the radius, and no claim that it estimates how wrong an expert is. No ordering claim about suppliers. No use of the word “first”. No claim that knowledge at distance at least one cannot harm. No claim of robustness under an estimated CPDAG. Retraction as a minimal correction set is cited rather than named as ours.

3 What we compare ourselves to, and where it enters the float

Requirement R3 is the one most easily faked. A related-work paragraph naming eight papers is not a comparison; a table in which a named prior procedure occupies a row and gets a number is. The list below is every piece of prior work that can occupy a row or a column of a Table 1 on this object, what it becomes when it does, and whether we have run it.

prior work	what it is	where it enters a Table 1	measured?
Perković et al., generalised adjustment criterion; Henckel et al., optimal set	the construction that turns a graph into the committed set	the estimator, identical in every row of every candidate; also the audit oracle	and it is a shared defect to fix: both sides’ validity checks must use this criterion rather than the back-door one, and the two disagree on a handful of committed sets
Guo and Perković, enumeration in an MPDAG (arXiv:2010.08611)	the interval over every adjustment set the class admits	the row that discards the knowledge, and the denominator of the calibrated width	every cell
Taeb, Guo and Henckel, model-oriented graph distances (arXiv:2511.10625)	the order and the distance the whole object is built on	the instrument that counts elementary graph revisions, one rung below ours	3,014 rows
structural Hamming distance	the standard graph-distance baseline	a column of the ranking table	synthetic arm only
number of asserted claims $ K $	the naive instrument an analyst already has	a rung of the certificate ladder and a column of the ranking table	both arms
treatment-to-set separation	a free structural rule read off the analyst’s graph	the cheapest rung of the certificate ladder	3,014 rows
leave-one-out and leave-two-out re-closure screen	not in the literature, and the first thing a competent reviewer proposes	a row of the price table and of the policy table	in the policy panel; – missing from the ranking table, and owed
Cinelli and Hazlett, omitted-variable bound	the sensitivity analysis the field actually runs	one row at its own oracle parameter, with the count of rows where it is defined	defined on 666 of 808 rows; covers 0.964 at 0.445 of the class width against 0.632 for ours on the same rows; on the 101 rows where the committed set contains a descendant of the treatment it covers 0.134
Rosenbaum; Masten and Poirier, breakdown frontiers	the conceptual analogue and the reporting form	the shape of the report table, not a computable baseline on this object	by construction

Two things follow. The first is that a candidate satisfies R3 only if a named prior procedure gets a number in the same float, not a mention in the caption; by that test the price table and the certificate ladder pass, the report table and the funnel do not. The second is that the Cinelli and Hazlett row is the only comparison against a sensitivity analysis the field runs, and it is favourable to them, because at their own oracle parameter their bound is narrower than our interval. Printing it is what makes the table an argument rather than an advertisement, and the honest reading is that it covers by assuming a model that a reversed orientation sits outside: on the rows where the committed set contains a descendant of the treatment, the same formula at the same oracle covers about one row in seven.

4 The design space, and the parts of it that are already dead

A Table 1 on this object is fixed by five choices. Naming them is what makes the candidate list finite rather than a brainstorm.

axis	the options that are live
the row	a complete procedure an analyst runs; an instrument that reads the graph; a knowledge condition; a named published analysis; a stratum of the population; a network
the column	the number of wrong claims; the sample size; the instrument; the safety statistic; the stratum; the depth
the cell	calibrated width with raw coverage; an over-certification count; a rank correlation; a radius; a share of a population; a named claim in the analyst’s words
the population	the semi-synthetic parameters on real structures; the four networks with fitted coefficients; the synthetic instance family; the author-declared queries; the learned-graph panel
the comparison	against a procedure, against an instrument, against a distance, against nothing

4.1 Combinations that are already dead

These were killed by measurement, not by taste. Rebuilding any of them costs a round and ends where it ended before.

the design	why it is dead
any table where the radius competes on discrimination	the expected rank of a reversed claim is $1 + (K - 1)(1 - \text{AUC})$, so a questions-to-repair column is a restatement of an area under a curve that already ties at 0.8417 against 0.8417. None of the discrimination tables survived the ideation run
median width as the width statistic	it is 1.00 in every cell of the coverage family
columns at $r = 0, 1, 2$ drawn as three populations	they are not nested, so the columns cannot be read across. The repair is one population per panel and columns indexed by the number of wrong claims
the radius as a width budget	widening over the whole class covers 0.9306 against 0.8835 for one revision. The claim was withdrawn in September and the calibrated metric is what answers the objection: after calibration the interval at one revision needs 0.851 of the class width
a coverage column and a width column, uncalibrated	discarding the knowledge covers 1.000 everywhere and passes any coverage gate
a tolerance band chosen so our row passes	the second repair the archetype adversaries demanded, after they found one passing by nine thousandths
an unpriced “verify by hand” column	a procedure in the table with no cost in the table
metrics as rows and procedures as columns	per-row dispersion becomes impossible, which is the seventh repair
the silent-failure grid	ranked first by the archetype run and refuted 0 of 3 by its own adversaries; one of its cells existed nowhere, and its rates mixed two denominators
a Table 1 whose rows are suppliers	ruled out on 12 September. Elicited knowledge is one error process in the methodology

4.2 What the measurements opened after the last design pass

The design runs ended on 9 September and the measurements landed on 12 and 13 September. Four candidates below could not have been written before that, and they are the reason this is not simply a re-reading of the old ranking: the certificate ladder aggregated over every knowledge condition (T2), the policy table with the refusal rate and the tie (T8), the report table with real bands and named breaking claims (T3), and the learned-graph scope panel (T7).

5 The candidates

Thirteen. Each is printed with the numbers that exist today; an unrun cell is marked and never left blank. Under each float is the file the numbers were read from, and after it a fixed block: what the table presents, the claims it verifies, the prior work that gets a number inside it, the objection it invites, and what it would cost to finish.

T1 The price of a wrong claim

Rows are complete procedures an analyst runs between receiving the graph and reporting an effect, with the same estimator in every row, ordered by how many revisions of the knowledge they tolerate. The columns are how many of the analyst’s claims are false.

procedure	$n = 1000$			$n = 20\,000$		
	$b = 0$	$b = 1$	$b = 2$	$b = 0$	$b = 1$	$b = 2$
report the estimate as committed	0.464 / .967	5.123 / .597	4.082 / .554	0.407 / .966	14.368 / .573	10.216 / .516
retract one claim far from the query	0.655 / .967	10.771 / .640	7.539 / .614	0.508 / .969	30.771 / .615	19.408 / .568
interval over one revision	0.957 / .982	0.851 / .980	4.234 / .781	0.922 / .980	0.758 / .981	14.031 / .757
interval over two revisions	0.984 / .984	0.980 / .984	0.847 / .984	0.971 / .983	0.965 / .984	0.801 / .981
discard the knowledge, whole class	1.000 / .984	1.000 / .985	1.000 / .988	1.000 / .984	1.000 / .985	1.000 / .985
informative problems (networks)	391 (23)	359 (23)	131 (9)	391 (23)	359 (23)	131 (9)

Draft caption. Width at 95 % calibrated coverage, relative to the interval over the whole equivalence class, with raw coverage beside it. Every row uses the same estimator and the same committed set; rows differ only in how many revisions of the asserted knowledge the reported interval tolerates. Columns are the number of the analyst’s claims that are false. Semi-synthetic parameters on published structures, 23 networks, network cluster bootstrap. The interval costs about twice the committed report when nothing is wrong, and one wrong claim is bought back at 85 % of the price of discarding the knowledge.

presents	the price of insurance in the currency the analyst reports in
verifies	C4 in full; C7 through the far-retraction row
prior work inside the float	the whole-class interval of Guo and Perković as the last row and as the denominator of every cell; the estimator is the generalised criterion
satisfies	R1 • R2 • R3 • R4 • R5 • R6 • (the interval loses at $b = 0$ by a factor of two) R7 • R8 •
the objection	the columns are a controlled corruption, not an error process anyone observed. The answer is T1b
status	every cell measured
cost to finish	none

T1b The same table, with the error process measured instead of imposed

Identical rows. The columns are replaced by the number of false claims the elicited knowledge actually contained in the treatment’s component.

procedure	$n = 1000$			$n = 20\,000$		
	none false	one false	two or more	none false	one false	two or more
report the estimate as committed	0.456 / .961	6.921 / .600	∞ / .242	0.365 / .964	15.639 / .552	∞ / .202
retract one claim far from the query	0.622 / .962	7.680 / .664	∞ / .296	0.443 / .964	19.828 / .605	∞ / .229
interval over one revision	0.977 / .981	0.867 / .976	6.011 / .531	0.963 / .980	0.854 / .982	19.711 / .499
interval over two revisions	0.997 / .982	0.979 / .982	7.671 / .636	0.995 / .981	0.955 / .986	28.746 / .607
discard the knowledge, whole class	1.000 / .982	1.000 / .983	1.000 / .993	1.000 / .981	1.000 / .988	1.000 / .986
informative problems (networks)	450 (20)	208 (13)	124 (7)	450 (20)	208 (13)	124 (7)

presents	the same price, under errors that were made rather than injected
verifies	C4, and C10 silently, without a sentence whose subject is a supplier
the objection	two or more false claims defeat every procedure short of the whole class, and the committed report’s width ratio is unbounded there. That is a finding and it has to be printed, but it is also the cell a hostile reader quotes
status	every cell measured
against T1	more honest and less controlled. Pooling suppliers keeps R8 satisfied; the column definition depends on an audit that reads the truth, which T1’s does not

T2 The certificate ladder: which instruments are safe to trust

Rows are instruments, ordered by cost, all of them computed from what the analyst has. Columns say how often each certifies an adjustment set that is invalid at the truth.

instrument	reads	certifies a set invalid at the truth		detection	false alarm
		rows	share		
number of claims that orient an edge	the claims	791	.262	.123	.046
number of claims asserted	the claims	729	.242	.192	.065
separation of treatment and set	the graph	15	.005	.983	.505
radius in elementary graph revisions	graph, claims, set	39	.013	.957	.203
radius in the analyst’s own claims	graph, claims, set	0	.000	1.000	.277

Draft caption. Five instruments an analyst can compute without the data-generating graph, ordered by cost. An instrument over-certifies when it declares an adjustment set safe at a depth at which the set is in fact invalid. Over 3,014 queries on 27 published network structures, under eleven knowledge conditions ranging from chance orientations to truthful ones, a certificate denominated in elementary graph revisions over-certifies on 39 queries and a certificate denominated in the analyst’s own claims on none. The zero follows from the criterion rather than from the sample: if the committed set is invalid at the truth and the analyst asserted w false claims, retracting those w claims leaves a graph the truth extends, so the claim radius is at most w whenever the search reaches that depth.

presents	the reason the radius is reported in claims and not in graph revisions
verifies	C3 in full, C2 in passing
prior work inside the float	the distance in elementary graph revisions is the model-oriented distance of Taeb et al. applied as an instrument; the two counts are the naive instruments and one of them is a column of the draft’s own ranking table
satisfies	R1 • R2 • R3 • R4 • (one safety statistic, with its two error kinds) R5 • R6 • (the last row has the worst false-alarm rate of the three usable instruments) R7 • R8 •
the objection	a zero that follows from a theorem is not evidence. The answer is the contrast: the same argument does not go through in the revision unit, and 39 rows show where it fails and: eleven conditions pooled into one denominator hides variation. The per-condition intervals exist and go in the appendix
status	every cell measured
cost to finish	none

T3 The report on named analyses

Rows are published analyses whose authors declared a treatment and an outcome. The table is laid out as a sensitivity report, with the interval first and the radius beside it.

analysis	query	K	true	estimate [95 %]	interval, 1 rev.	interval, 2 rev.	r_0	first breaking claim
Didelez 2010	HRT on TCI	1	−.586	−.586 [−.599, −.572]	[−.604, −.568]	[−.604, −.568]	never	none to depth 3
Shrier 2008	warm-up on injury	2	−1.351	−1.348 [−1.358, −1.339]	[−1.431, −1.228]	[−1.431, −1.228]	never	none to depth 3
Polzer 2012	tooth loss on mortality	5	+1.193	+1.190 [+1.183, +1.197]	[+1.128, +1.208]	[+1.128, +1.208]	never	none to depth 3
Sebastiani 2005	EDN1.3 on EDN11.7	7	−1.212	−1.209 [−1.219, −1.199]	[−1.234, −1.179]	[−1.234, −1.179]	never	none to depth 3
mediator (textbook)	X on Y	3	−2.603	−2.600 [−2.623, −2.576]	[−2.623, −1.109]	[−2.623, +0.000]	2	Z→I and Z→X
ecoli70	asnA on aceB	10	+547	+542 [+511, +572]	[−.011, +.572]	[−.011, +.572]	1	asnA→icdA
ecoli70	pspB on nmpC	10	−108	−100 [−.118, −.082]	[−.531, +.026]	[−.531, +.026]	1	pspB→pspA
arth150	100 on 245	29	+1.535	+1.534 [+1.527, +1.542]	[+0.000, +1.545]	[+0.000, +1.545]	1	100→245
magic-niab	G311 on G43	8	−.255	−.251 [−.259, −.244]	[−.260, +0.000]	[−.260, +0.000]	1	G1853→G311
magic-irri	G1378 on SUB	10	−.035	−.031 [−.044, −.017]	[−.048, +.021]	[−.048, +.021]	1	G1378→G1311

Draft caption. A sensitivity report for ten analyses. The interval is over every knowledge state within the stated number of revisions of the analyst’s claims; the null radius is the number of revisions after which the interval contains zero, in the sense Rosenbaum’s value has for unmeasured confounding. On five of the ten, no retraction to depth three moves the conclusion. On the other five a single claim does, and it is named.

presents	the paper’s framing made concrete, in the reporting form the field already reads
verifies	C1, C2; and C8 inside itself, because the declared queries are the robust ones
prior work inside the float	none gets a number. The form is Rosenbaum’s and Masten and Poirier’s, and the analyses are other people’s
satisfies	R1 • R2 • R3 • R4 • R5 • R6 • R7 • R8 •
the objection	instance selection is the whole attack surface, and the eight-of-nine robust declared queries sit uncomfortably against the sampled population where almost everything breaks at the first retraction. Any inclusion rule has to be stated as a rule with the exclusions listed
status	measured. Three declared queries are excluded for stated structural reasons
verdict	this is Table 2, not Table 1. It has no comparison in it

T4 The ranking table, with the baseline it is missing

The draft’s current Table 1, with one column added. Rows are strata of the corruption design; columns are rank correlations between an instrument and how long the committed set survives.

stratum	n	$\tau_b(r_{\text{val}})$	$\tau_b(\text{SHD})$	$\tau_b(K)$	$\tau_b(\text{screen depth})$
flip, coverage 0.5, base wrongness 0.00	239	.634 [.533, .728]	.204 [.105, .300]	.184 [.076, .289]	unrun
flip, coverage 0.5, base wrongness 0.10	240	.334 [.206, .453]	−.000 [−.111, .107]	−.004 [−.114, .105]	unrun
flip, coverage 0.5, base wrongness 0.25	220	.279 [.153, .397]	.032 [−.063, .127]	.131 [.000, .256]	unrun
flip, coverage 1.0, base wrongness 0.00	229	−.091 [−.204, .030]	undefined	.021 [−.062, .103]	unrun
flip, coverage 1.0, base wrongness 0.10	226	.474 [.393, .548]	.263 [.160, .367]	.261 [.171, .349]	unrun
flip, coverage 1.0, base wrongness 0.25	133	.355 [.223, .476]	.582 [.488, .666]	.557 [.473, .635]	unrun
tiered, 2 tiers	190	.791 [.763, .814]	.361 [.287, .434]	−.423 [−.516, −.327]	unrun
tiered, 3 tiers	238	.756 [.728, .781]	.169 [.079, .257]	−.382 [−.459, −.303]	unrun
tiered, 4 tiers	240	.722 [.690, .750]	.195 [.103, .286]	−.377 [−.453, −.300]	unrun
flip, coverage 0.5, base wrongness 0.25, 1000 draws	220	.032 [−.112, .168]	−.170 [−.260, −.080]	−.200 [−.334, −.068]	unrun

presents	that the radius orders instances by survival better than the distances an analyst would otherwise use
verifies	C11
prior work inside the float	structural Hamming distance and the claim count, both with numbers. This is the strongest R3 score of any candidate
satisfies	R1 ◦ (a reader must already accept the survival area) R2 ◦ R3 • R4 • R5 • R6 • (the radius loses to both baselines in the sixth stratum, and one stratum is a confirmed null) R7 – (this is a discrimination table, which the ideation run killed for our claim set and which is legitimate for his) R8 •
the objection	the missing column. A reviewer who thinks of the leave-one-out screen asks why it is not beside the two weak baselines, and our own policy panel says it ties the radius at matched depth. Adding the column is honest and reduces the claim
status	three columns measured on synthetic structure only; the real-instance evaluation is what our protocol run supplies; the fourth column is unrun
cost to finish	the screen column on the synthetic arm is cheap, because the corruption states are already stored. Estimate half a day
verdict	this is the ranking section’s table. As Table 1 of a merged paper it asks the reader to hold a survival area in their head before they have seen what breaks

T5 The decision table

One row per query, in the analyst’s own variable names, saying which single claim the conclusion hangs on. The columns left of the rule are computed from observables; the two on the right read the truth and are the audit.

network	query	radius	holds unless this is retracted	free rule	named claim false	set at truth
Didelez 2010	Occ on HRT	1	Age → Occ	1	no	valid
Didelez 2010	Smo on HRT	1	Age → Occ	2	no	valid
Acid 1996	x1 on x10	1	x4 → x1	1	yes	invalid
Polzer 2012	Caries on mortality	1	Caries → Psychosocial	undefined	yes	invalid
Kampen 2014	ALN on PER	2	AFF → ALN and ALN → PER	1	no	valid
Kampen 2014	AIS on DET	3	ALN→SAN, APA→SAN, SAN→AIS	2	yes	invalid
Schopf 2010	PA on T2DM	1	A → PA	1	no	valid

229 certifiable queries; 157 carry an exact certificate; the breaking set is a single claim on 141 of those 157.

In the audit the named claim was false on 68 of 157. Where it was false the committed set was invalid at the truth on 58 of 68; where it was true, the set was valid on 85 of 89. The sentence points at the claim that in fact broke the set on 58 of the 62 invalid rows.

Draft caption. What the diagnostic returns to an analyst. For each query the certificate names the smallest set of the analyst’s own claims whose retraction invalidates the committed adjustment set; on 141 of 157 queries with an exact certificate it is a single claim, so the report is a sentence. The two right-hand columns are the audit, computed after the rows were frozen: whether the named claim was in fact false, and whether the committed set was in fact invalid.

presents	the output of the method as a practitioner receives it, with its own audit attached
verifies	C2, and the well-posedness of the report
prior work inside the float	the free rule column only
satisfies	R1 • (the most legible candidate by a distance) R2 ◦ R3 – R4 – (no metric at all) R5 • R6 – R7 • R8 •
the objection	it is an illustration, not a result. Nothing is compared and nothing is priced
status	measured, including the audit
verdict	this is the figure or the first table of the results section, and it is the best single artefact for the introduction. It is not Table 1 because it has no comparison

T6 The applicability funnel

Rows are tiers of the substrate. Columns say what survives from a parsed network file to a query the method can certify.

tier	networks parsed / framed / swept	open edges	candidate pairs	sampled	certifiable	effective n
textbook diagrams	3 / 2 / 2	19	252	32	13	32.0
published applied graphs	24 / 19 / 13	615	76,565	380	91	51.0
benchmark networks	8 / 8 / 8	41	482	156	99	41.6
large gene networks	4 / 4 / 4	92	2,833	80	26	19.5
total	39 / 33 / 27	767	80,132	648	229	

presents	where the method applies, and how much of the corpus it reaches
verifies	C8
satisfies	R1 ● R2 ○ R3 – R4 – R5 ● R6 – R7 ● R8 ●
the objection	none. It is the answer to an objection
verdict	a funnel as Table 1 sells a scope limit as the headline. This belongs early in the evaluation section, where it disarms the applicability attack before a reviewer assembles it

T7 The scope panel: what happens when the graph is estimated too

The same instruments, on four networks whose CPDAG was learned from data the true graph generated rather than read from the true graph.

knowledge condition	graph from the truth		graph learned from data		
	certifiable	set invalid at truth	certifiable	set invalid at truth	over-certifies
language model, real names	26	.269	38	.737	14
chance, matched on the same edges	66	.409	116	.733	43
no knowledge asserted	11	.000	19	.737	14
truthful orientations	18	.000	34	.676	23
truthful orientation of every open edge	37	.000	37	.595	22

presents	the boundary of the licensed claim, with a number on it
verifies	C9
the finding to read	even truthful knowledge over-certifies once the graph itself is wrong, on 23 of 34 rows. Discovery error is not a knowledge error and no retraction of the claims repairs it
satisfies	R1 ● R2 ● (of the scope claim) R3 – R4 – R5 ● R6 ● R7 ● R8 ●
verdict	a results paragraph and a limitation. Printing it as Table 1 would be an unusual act of self-harm, and not printing it at all would be the defect the draft’s discussion already admits it has

T8 The policy table: what the certificate is worth against a free screen

Rows are decision rules an analyst could follow. The first column is what the rule costs when the knowledge is entirely correct; the rest is what it buys when it is not.

policy	refuses when nothing is wrong	one wrong claim		two wrong claims	
		coverage	width	coverage	width
report as committed	.000	.597	5.123	.554	4.082
leave one claim out, re-close, widen if the set changes	.910	.954	0.691	.753	5.967
certify at depth one, widen if refused	.910	.979	0.715	.776	4.640
certify at depth two, widen if refused	.955	.984	0.980	.986	0.791
leave one or two out, re-close, widen if the set changes	.963	–	–	–	–

The two comparisons this table exists to make, both fixed before they were computed. Certifying at depth two beats the leave-one-out screen by 0.305 [0.182, 0.410] in coverage at the population and by 0.256 [0.149, 0.346] at a sample of one thousand. It ties the leave-one-or-two-out screen at 0.000 [0.000, 0.000] in coverage and –0.002 in width, and the two rules decide differently on 5 of 754 rows. The separation the design asked for exists because the baseline it named stops at depth one; matched on depth, the free screen costs the same kind of graph operation and returns the same answer.

presents	the honest position of the method against the baseline a reviewer invents
verifies	C5 and C6
prior work inside the float	the screen. It is absent from the literature and it is the comparison that matters most
satisfies	R1 ● R2 ● R3 ○ R4 ● R5 ● R6 ● (the refusal column is a defeat and it is the first column) R7 ● R8 ●
the objection	printing this invites the reader to conclude that the radius is unnecessary. The answer is the depth, the named witness and the certificate that no failure lies nearer, none of which a screen returns. The screen itself is safe, and the audit says so: on all 2,212 rows, a refusal at depth r is matched by a screen that triggers at depth at most r , with no exception. What the ladder of T2 shows to be unsafe is the graph-revision unit, which is a different object
status	measured
verdict	a strong and unusual Table 1: it leads with what the method does not buy. As a second table it is the paper’s credibility

T9 The distribution of the quantity

Rows are knowledge conditions, columns are the value the certificate returned.

knowledge condition	certifiable	$r = 1$	$r = 2$	$r = 3$	beyond 3	no retraction breaks it
language model, real names	229	141	8	8	44	27
same family, larger model	285	190	15	0	50	16
chance, matched on the same edges	649	395	32	12	129	81
truthful orientations	209	128	7	1	45	27
truthful orientation of every open edge	294	226	0	0	46	19

presents	what the quantity actually looks like on real structure
verifies	nothing by itself; it disarms the degeneracy objection by printing it
the finding to read	the conclusion rests on a single claim on 62 % of certifiable queries under the primary condition, and on 77 % under a fully oriented graph. The certificate’s measured range on real structure is one, two or three, and every sentence about a larger radius is a sentence about a hop count
satisfies	R1 • R2 – R3 – R4 – R5 • R6 • R7 • R8 ◦
verdict	an appendix table, or a bar in a figure. It is the honest answer to “your quantity is degenerate” and it cannot be Table 1 because it verifies no claim

T10 The supplier ledger (listed for completeness; ruled out)

Rows are knowledge conditions, columns are what the audit found once the truth was revealed.

condition	certifiable	claims false	set invalid at truth	radius over-certifies	hop over-certifies
7B, real names	229	.329 [.184, .494]	.288 [.155, .435]	0	4
7B, second presentation order	206	.323 [.190, .483]	.277 [.149, .426]	0	3
second family, 8B	243	.285 [.165, .420]	.230 [.130, .344]	0	5
7B, scrambled names	197	.394 [.296, .484]	.340 [.211, .467]	0	7
same family, 72B	285	.278 [.163, .405]	.298 [.171, .431]	0	1
chance, matched on the model’s edges	649	.479 [.398, .559]	.378 [.295, .460]	0	10
truthful orientations	209	.000	.000	0	0

verifies	C10
why it is ruled out	the ruling of 12 September. Elicited knowledge is one error process in the methodology and never the headline, the title or Table 1. It is listed here because the brief asked for every possible Table 1, and because it is the most quotable table we have, which is exactly the risk
where it goes	the evaluation section as one paragraph and two rows, and the appendix in full
the finding worth keeping	a larger model asserts more claims and gets more of them right, and the share of committed sets that are invalid at the truth does not move: 0.298 against 0.288. Accuracy and validity are not the same purchase

T11 Two panels: is it safe, and what does it cost (recommended)

One float, two panels, two populations named separately and never pooled. Panel A is the certificate ladder; Panel B is the price table with the screen row added.

Panel A. Over-certification, 3,014 queries on 27 published structures, eleven knowledge conditions

instrument	certifies an invalid set	share	detection	false alarm
number of claims asserted	729	.242	.192	.065
separation of treatment and set	15	.005	.983	.505
radius in elementary graph revisions	39	.013	.957	.203
radius in the analyst’s own claims	0	.000	1.000	.277

Panel B. Calibrated width and coverage, 391 problems on 23 of the same structures, $n = 1000$

procedure	no wrong claim	one wrong claim	two wrong claims
report the estimate as committed	0.464 / .967	5.123 / .597	4.082 / .554
leave one claim out, re-close, widen if it changes	0.932 / .981	0.691 / .954	5.967 / .753
interval over one revision of the claims	0.957 / .982	0.851 / .980	4.234 / .781
interval over two revisions of the claims	0.984 / .984	0.980 / .984	0.847 / .984
discard the knowledge, interval over the whole class	1.000 / .984	1.000 / .985	1.000 / .988
informative problems (networks)	391 (23)	359 (23)	131 (9)

Draft caption. Panel A: an instrument over-certifies when it declares an adjustment set safe at a depth at which the set is invalid in the data-generating graph. Over 3,014 queries on 27 published structures, under knowledge ranging from chance orientations to truthful ones, a certificate counting elementary graph revisions over-certifies on 39 queries and one counting the analyst’s own claims on none. Panel B: width at 95 % calibrated coverage relative to the interval over the whole equivalence class, with raw coverage beside it, over the problems on which the class leaves the effect open. The insurance costs about twice the committed width when nothing is wrong, and one wrong claim is bought back at 85 % of the price of discarding the knowledge; the leave-one-out screen reaches the same decision as the certificate at this depth.

presents	the two questions a reader has after the abstract, in the order they have them: can I trust the certificate, and what does acting on it cost
verifies	C3 and C4 in full; C5 through the screen row, which sits beside the interval it matches; C6 through the refusal number printed under the float
prior work inside the float	the model-oriented distance as the revision-unit instrument, the claim count, the whole-class interval of Guo and Perković, and the screen. Four named comparisons with numbers
satisfies	R1 ● R2 ● R3 ● R4 ● R5 ● (two populations, separately named) R6 ● (the method loses the first column of Panel B and has the second-worst false-alarm rate in Panel A) R7 ● R8 ●
the objection	two panels is a lot for one float on a nine-page limit, and the two populations have different denominators, so the eye wants to read across and must not. The caption has to say so
status	every cell measured
cost to finish	none. The optional sensitivity row is a further half day

T12 Three panels, for the merged paper

T11 with the ranking panel above it, so the draft’s own result leads and ours supports it.

panel	content
A, does it rank	four strata of the corruption design, the rank correlation of the radius against structural Hamming distance, the claim count, and the screen depth. Three columns exist; the fourth is unrun
B, is it safe	the certificate ladder, five instruments, 3,014 queries
C, what does it cost	the price table, five procedures, three columns

presents	the whole argument in one float: the quantity orders instances, the unit is safe, and the report has a price
verifies	C3, C4, C11
satisfies	R1 ○ R2 ● R3 ● R4 – (three metrics) R5 ● R6 ● R7 – R8 ●
the objection	three panels on three populations with three metrics is a results section wearing a float. At a nine-page limit it eats most of a page and a reviewer reads none of it carefully
verdict	the honest merge, and the wrong shape. Panels B and C are T11; panel A is the ranking section’s own table

T13 The unit disagreement

Rows are instruments, columns are how their reading compares with the depth measured in the analyst’s claims.

instrument	reads below	agrees	reads one above	reads two above	largest gap
separation of treatment and set	14	136	7	0	2
radius in elementary graph revisions	1	135	19	2	2

presents	the unit mismatch as an agreement table rather than a safety table
verifies	C3 and C7
satisfies	R1 ● R2 ○ R3 ○ R4 ● R5 ● R6 – R7 ● R8 ●
the objection	agreement is not safety, and 21 rows above out of 157 reads as a small disagreement until the reader is told that every one of them is a query certified safe at a depth at which it is not
verdict	a panel of T2, or a sentence in the locality section

6 The scorecard

The columns are the three requirements from your brief and the six constraints already paid for, plus how much of the table exists today and how it sits against the draft that arrived this morning. A dot is satisfied, a circle is partly, a dash is not.

	candidate	R1	R2	R3	R4	R5	R6	R7	R8	measured	fits the draft	natural place
T1	price of a wrong claim	●	●	●	●	●	●	●	●	all	●	Table 1 or 2
T1b	the same, errors measured	●	●	●	●	●	●	●	○	all	●	a variant of T1
T2	the certificate ladder	●	●	●	●	●	●	●	●	all	●	Table 1
T3	report on named analyses	●	○	—	●	○	—	●	●	all	●	Table 2
T4	ranking, with the baseline	○	○	●	●	●	●	—	●	3 of 4 cols	●	the ranking section
T5	the decision table	●	○	—	—	●	—	●	●	all	●	Figure 1 or first result
T6	the applicability funnel	●	○	—	—	●	—	●	●	all	●	evaluation setup
T7	the scope panel	●	●	—	—	●	●	●	●	all	●	limitations
T8	the policy table	●	●	○	●	●	●	●	●	all	○	Table 2 or 3
T9	the distribution	●	—	—	—	●	●	●	○	all	●	appendix
T10	the supplier ledger	●	○	—	—	●	●	●	—	all	○	appendix
T11	two panels, safe and costly	●	●	●	●	●	●	●	●	all	●	Table 1
T12	three panels, merged	○	●	●	—	●	●	—	●	most	●	too large
T13	the unit disagreement	●	○	○	●	●	—	●	●	all	●	a panel of T2

6.1 Which claim each candidate would actually verify

A table verifies a claim when a reviewer could reject the claim by reading the table and finding the wrong number in it. By that test most of the candidates verify nothing, and this is the column that separates them.

	C1	C2	C3	C4	C5	C6	C7	C8	C9	C10	C11
	object	computable	unit	price	tie	refusal	locality	reach	scope	supplier	ranking
T1				●			○				
T1b				●						○	
T2		○	●				○				
T3	●	●						○			
T4											●
T5		●									
T6								●			
T7									●		
T8					●	●					
T9		○									
T10										●	
T11		○	●	●	○	○	○				
T12		○	●	●	○	○					●
T13			○				●				

6.2 The recommendation

Table 1 is T11. Panel A is the certificate ladder, panel B is the price table with the screen row and the refusal column. Every cell exists. It carries the two claims that are ours and measured, C3 and C4, and it gives a number to four named comparisons inside the float: the model-oriented distance as the revision-unit instrument, the claim count, the whole-class interval, and the free screen. It is also the only candidate that loses somewhere a reviewer can see, in two different places, which is what R6 asks for.

The rest of the paper’s tables follow from that choice. T3 becomes Table 2, the report on named analyses, which is where the framing of C1 becomes concrete and where the Rosenbaum form pays. T8 becomes Table 3 if there is room, and the tie against the depth-matched screen goes into the results prose if there is not; it is the paper’s credibility and it must appear somewhere with its numbers. T4 stays where it is, in the draft’s ranking section, with the screen column added. T6 and T7 are the evaluation setup and the limitation. T5 is the best candidate for the first figure, because it is the only artefact that shows what an analyst receives. T9, T10 and T13 are appendix material.

Two alternatives are defensible and neither is better. If the merged paper keeps the diagnostic framing rather than the pricing one, Table 1 is T2 alone and T1 becomes Table 2; the argument then runs through validity and the interval arrives later, which costs the practitioner sentence its place on page one. If the ranking result is judged the strongest thing either side has, Table 1 is T4 with its fourth column and everything else moves down; that is the shape of the draft as it stands, and it asks a reader to accept a survival area before they have seen anything break.

6.3 What is blocking, and what is merely owed

Before the abstract on 18 September.

The novelty search for C1 has never been run in the wording the claim uses. The register entry says so explicitly and refuses to carry a verification date. The sensitivity-analysis lineage the claim excludes was never searched, and the framing rests on a statement about a paper nobody on this side has read in full. A claim of this shape, submitted unsearched, is the exact failure that cost the tutorial last July. Either the search runs, or C1 is narrowed to what was searched and the abstract says the narrower thing. The author list is the other lock. It is written nowhere, and one collaborator has been on every message since 28 August without being asked.

owed	why	cost
the screen column of T4	a reviewer who knows the leave-one-out screen asks for it beside the two weak baselines, and our own policy panel says it ties at matched depth. Adding it is honest and it reduces the ranking claim	half a day, the corruption states are stored
one validity oracle for both sides	the audit must use the generalised adjustment criterion rather than the back-door criterion; the two disagree on a handful of committed sets, and the draft and the ledger currently do not share one	hours
the anti-exchange caveat replaced by the number	the draft carries it as an unproved property; the audit gives 5,000 draws, zero violations and a 0.06% upper bound	minutes, it is written
the sensitivity row of T1	the only comparison against a sensitivity analysis the field runs, favourable to them at their own oracle	half a day
the corrupted-knowledge calibration on the frame	being rerun after the seed fix; T1's controlled columns depend on it	running

6.4 One decision that is not mine

The framing conflict is real and it is yours to settle. The draft's abstract already says the radius prices structural risk, which is the pricing move; the quantity it prices is validity alone. Ours leads with the interval and reports the radius beside it. T11 is built for the second reading and degrades gracefully into the first, because panel A stands on its own. If you take the draft's framing whole, the recommendation changes to T2 and the two panels separate.